Predictive Advantage, Regime Dependence, and Direction-Specific Causal Control in Kinetic Ising Dynamics

Kunal Bhatia[1]

[1]*Independent Researcher, Heidelberg, Germany*

Abstract

We ask whether predictive sufficiency implies causal access in the two-dimensional kinetic Ising model, and whether predictive ranking is stable across dynamical regimes. The answer to both is no. The central result is a predictor-handle dissociation: a coarse macro-state and the fine champion are statistically indistinguishable as predictors of future domain-wall change, yet only the fine structure yields a direction-specific causal handle. A regime-crossover result earns the regime-dependence qualifier and blocks the objection that fine descriptions win simply by having more bits. Four empirical axes support these claims. (**Axis 1, organism comparison**) Against a block-variance coarse rival that achieves $R^2 \approx 0$ for ΔW in the ordered phase (a near-null baseline), the fine champion $E_{\geq 3} + J$ achieves $\Delta R^2 = +0.391$ (95% CI [+0.367, +0.418]). Against the stronger coarse rival $[W_0, B_{\text{mag}}]$, the gap narrows to $\Delta R^2 = 0.004$ (CI includes zero): the two descriptions are statistically indistinguishable as predictors. (**Axis 2, regime crossover**) This advantage reverses sign in the disordered phase (+0.178 ordered, −0.071 disordered), showing that the fine-over-coarse advantage is regime-dependent. (**Axis 3, target specificity**) Under a fine-over-coarse ΔR^2 criterion, future domain-wall change ΔW is the target most specifically favoured by the fine description ($\Delta R^2 = +0.392$); this is not merely the most predictable target in absolute terms. (**Axis 2, regime crossover**) includes a near-critical probe at $T = 2.20$, which finds $\Delta R^2 \approx -0.013$ (CI spans zero), placing the crossover between $T = 1.80$ and $T = 2.20$. (**Axis 4, causal control**) Biasing the Glauber generator toward wall-reducing flips at exposed sites accelerates future wall removal (aligned arm removes more walls than baseline), while biasing against those flips retards it; this direction-specific sign flip holds across all active perturbation strengths (4/4 sign flips on non-null ε values). A class-level control (budget-matched random-site arm with the same directional logic) shows that the operative causal mechanism is the per-site directional conditioning—not $d_c = 4$ class identity specifically—and that global aggregate descriptions cannot provide this per-site handle. (**Dissociation**) The aggregate coarse macro-state $[W_0, B_{\text{mag}}]$ ($R^2 = 0.389$) and the fine champion $E_{\geq 3} + J$ ($R^2 \approx 0.391$) are statistically indistinguishable as predictors of ΔW, yet only per-site fine structure enables a direction-specific causal handle: global aggregate descriptions carry no per-site directional information and cannot be used to implement the generator bias of Axis 4. Two descriptions with equivalent predictive power differ entirely in their capacity to support per-site causal intervention.

PACS numbers: 05.10.Ln, 05.50.+q, 75.10.Hk, 64.60.Cn

I. INTRODUCTION

A common working assumption in statistical physics is that different descriptions of the same system may differ in predictive usefulness—but which description is best, and whether that ranking is stable across regimes and tasks, is not always obvious. This paper studies how predictive and causal usefulness depend on the choice of description in a simple, well-understood dynamical system.

The central question is whether predictive sufficiency implies causal access: can a description that accurately forecasts future domain-wall change be used to controllably change it? We find the answer is no, and that the dissociation is not an artefact of a particular baseline choice: even the strongest available coarse rival $[W_0, B_{\mathrm{mag}}]$ matches the fine champion predictively ($\Delta R^2 = 0.004$, CI includes zero), yet only the fine structure provides a direction-specific causal handle. The crossover result earns the regime-dependence qualifier: the fine-over-coarse predictive advantage reverses sign from ordered (+0.178) to disordered (−0.071 phase), blocking the objection that fine descriptions simply win by having more bits. The fine champion is $E_{\geq 3} + J$; the target most specifically favoured by fine descriptions is ΔW; and the causal sign-flip holds across all 4/4 active perturbation strengths (the $\varepsilon = 0$ null produces a gap consistent with zero, as expected).

a. Why the Ising model. The 2D kinetic Ising model with Glauber dynamics [4, 7] is one of the few systems for which domain-wall coarsening, defect densities, and critical exponents are all precisely characterised [1]. Domain walls provide a structurally meaningful target family: interfaces are the elementary reorganisation objects under coarsening, and their evolution is driven by specific local geometric features of the current configuration. An isolated minority spin surrounded by four unlike neighbours—a site with exposure degree $d_c = 4$—contributes four domain-wall bonds, is energetically costly, and is a natural structural precursor of upcoming interface simplification. This geometry makes the Ising model a clean and interpretable testbed.

b. Why description choice matters. Finer descriptions always contain more bits, but more bits do not automatically produce better predictions for a specific target. The predictive signal for interface-level change lies in sparse, geometrically special objects—exposed tips,

junction bends—that coarse block averages discard. In the disordered phase no such sparse structure exists; coarse summaries recover the predictive advantage. And predictive sufficiency does not confer causal access: as we demonstrate, the variable that best forecasts ΔW is not the variable whose structure can be used to controllably change ΔW.

c. Paper organisation. This paper establishes four empirical results and one dissociation result: **(1)** organism comparison in the ordered phase; **(2)** regime crossover demonstrating regime-dependence; **(3)** target specificity under a fine-over-coarse criterion; **(4)** direction-specific causal control via generator alignment; and **(D)** predictor-handle dissociation. Section II reviews related work. Section III presents the model, observables, and causal protocol. Section IV reports results with figures and tables. Sections V–VI discuss implications and conclude.

II. RELATED BACKGROUND

a. Coarse-graining and renormalisation. The renormalisation group [9, 12] is the canonical framework for relating descriptions at different scales. In that tradition, coarse-graining maps fine configurations to block variables, and RG fixed points characterise universality classes. Our approach is complementary: we ask empirically, for a given target and horizon, which description achieves higher out-of-sample predictive accuracy. Israeli and Goldenfeld [8] showed that cellular automaton coarse-graining can produce exact reduced models, illustrating the system-specificity of which description is "best."

b. Prediction versus causation. The distinction between prediction and intervention is formalised in the do-calculus of Pearl [10] and the interventionist philosophy of Woodward [13]. A variable may be a strong predictor while being causally inert if it is correlated with the true causal driver. Although the principle that predictive sufficiency neither implies nor is implied by causal identifiability follows from the do-calculus in full generality, the question of *which specific physical descriptions* possess causal leverage — and whether this depends on the dynamical regime, the target, and the structural class targeted — is empirical and not answered by the theory alone. Our contribution is to demonstrate, quantify, and characterise the dissociation within the kinetic Ising model, providing a tractable physical testbed where both predictive accuracy and causal leverage can be measured simultaneously under controlled experimental designs (including a class-level structural control).

c. Learned coarse-graining and description selection. A growing literature uses machine learning to identify optimal coarse-grained variables for molecular and lattice systems [14, 15]. These approaches optimise predictive accuracy of a chosen target, raising the question of whether accurate predictors also serve as intervention handles. Our work addresses this question directly in a physical model: equivalent predictive power (in the $[W_0, B_{\mathrm{mag}}]$ vs. $E_{\geq 3} + J$ comparison) does not imply equivalent causal leverage.

d. Structural defects in coarsening. The role of domain-wall defects in Ising coarsening is well studied [1, 11]. Isolated minority spins and junction-type sites are energetically unstable and drive local coarsening events. We quantify how well such defects predict future macroscopic wall change and test whether they serve as intervention targets.

e. Macro-micro comparisons in prediction and causation. The causal emergence programme [5, 6] asks whether macro descriptions can have more causal organisation than micro ones, measured by effective information. Our approach differs: we use standard regression criteria for predictive accuracy and direct intervention experiments for causal leverage, without assuming any universal scale ranking.

III. MODEL AND METHODS

A. The 2D kinetic Ising model

We study the ferromagnetic Ising model on a 64×64 square lattice with periodic boundary conditions. Each site i carries spin $\sigma_i \in \{-1, +1\}$. Dynamics are governed by single-spin-flip Glauber updates [4]:

$$p(\sigma_i \to -\sigma_i) = \frac{1}{1 + \exp\big(2\beta h_i \sigma_i\big)}, \tag{1}$$

where $h_i = \sum_{\langle j \rangle} \sigma_j$ and $\beta = 1/T$. A checkerboard update scheme avoids spatial bias. The 2D Ising critical temperature is $T_c \approx 2.269$.

Two anchor temperatures are used: ordered $T = 1.80$ (well inside the ferromagnetic phase) and disordered $T = 2.50$ (above T_c). Each configuration is equilibrated for 1000 sweeps before measurement. Predictive targets are measured at horizon $K = 20$ sweeps.

B. Structural observables

a. Domain-wall count. $W(\sigma) = \sum_{\langle i,j \rangle} \mathbf{1}[\sigma_i \neq \sigma_j]$. Future change: $\Delta W = W(t_0 + K) - W(t_0)$.

b. Exposed-site counts. Exposure degree: $d_c(i) = \sum_{\langle j \rangle} \mathbf{1}[\sigma_j \neq \sigma_i]$.

$$E(\sigma) = \sum_i \mathbf{1}[d_c(i) = 4] \quad \text{(isolated minority spins)}, \tag{2}$$

$$E_{\geq 3}(\sigma) = \sum_i \mathbf{1}[d_c(i) \geq 3] \quad \text{(highly exposed sites)}. \tag{3}$$

c. Junction count. $J(\sigma) = \sum_i \mathbf{1}[d_c(i) = 2 \text{ and non-collinear}]$. Junctions mark the bends and corners of wall segments.

d. Composite descriptors. We use *descriptor* (equivalently, *organism* in the causal emergence literature [5]) to denote any deterministic map $f : \{-1, +1\}^{L^2} \to \mathbb{R}^d$ used as a regression predictor. Fine descriptors resolve local structural topology; coarse descriptors summarise global or block-averaged properties. $E_{\geq 3} + J$ is the two-column matrix $[E_{\geq 3}, J]$; each column receives a separate regression coefficient. $E + J$ uses $[E, J]$. These are not scalar sums.

e. Coarse aggregate (dissociation rival). The aggregate coarse macro-state used in the predictor-handle dissociation is $[W_0, B_{\text{mag}}]$, where

$$B_{\text{mag}}(\sigma) = \frac{1}{N} \Big| \sum_i \sigma_i \Big|, \quad N = L^2, \tag{4}$$

is the bulk magnetisation magnitude and $W_0 = W(\sigma(t_0))$ is the initial wall count. Together, $[W_0, B_{\text{mag}}]$ captures aggregate global information about configuration size and ordering without resolving interface topology.

f. Coarse block-variance rival.

$$\phi_{\text{var8}}(\sigma) = \frac{1}{n_b} \sum_{b=1}^{n_b} \left(|\bar{\sigma}_b| - \tfrac{1}{n_b} \sum_{b'} |\bar{\sigma}_{b'}| \right)^2, \tag{5}$$

the variance of block-averaged order parameter across 8×8 blocks.

C. Predictive evaluation

Ridge regression with penalty $\lambda = 1.0$, 5-fold cross-validation, $N = 3000$ configurations. A sensitivity analysis over $\lambda \in \{0.01, 0.1, 1.0, 10.0\}$ shows R^2 variation below 0.005 for all

comparisons reported here; conclusions are not sensitive to the regularisation choice. Primary statistic:

$$\Delta R^2 = R^2(\text{fine descriptor}) - R^2(\phi_{\text{var8}}). \tag{6}$$

Confidence intervals use pairwise bootstrap resampling ($B = 1000$) with the same configuration set for fine and coarse predictions.

a. Reproducibility. Stochastic dynamics are implemented via Numba-compiled functions; the Numba random-number generator is seeded independently from the worker-level seed to ensure cross-run reproducibility. Reported confidence intervals are stable to this level of control.

D. Causal protocol: aligned and misaligned generator bias

Rather than forcing state changes at t_0, we modify the Glauber transition probability throughout the K-step evolution. At each exposed site i ($d_c(i) = 4$), the local wall-change contribution of flipping σ_i is

$$\delta W_i = \sum_{\langle j \rangle} \big[\mathbf{1}[\sigma_j \neq -\sigma_i] - \mathbf{1}[\sigma_j \neq \sigma_i]\big]. \tag{7}$$

For a fully exposed site, $\delta W_i < 0$.

Two arms are defined:

Aligned:: $\delta W_i < 0$ transitions promoted ($p \to p + \varepsilon$); $\delta W_i > 0$ suppressed ($p \to p - \varepsilon$).

Misaligned:: Opposite: $\delta W_i < 0$ suppressed, $\delta W_i > 0$ promoted.

Both share initial configuration $\sigma(t_0)$ and are compared to an unperturbed baseline arm. The intervention identifies currently exposed sites ($d_c = 4$) at each of the K sweeps, not from a mask fixed at t_0; this ensures the structural targeting tracks the evolving configuration throughout the K-step evolution. Primary output:

$$\text{gap} = \Delta W_{\text{baseline}} - \Delta W_{\text{arm}}, \tag{8}$$

positive meaning the arm achieves greater wall reduction than baseline (i.e. aligned arm removes more walls, accelerating coarsening). The *sign-flip discriminator* requires aligned gap > 0 and misaligned gap < 0, both with CIs excluding zero. Five values of ε are tested: $\varepsilon = 0$ (null control, gap should be near zero with CI spanning zero) and $\varepsilon \in \{0.02, 0.05, 0.10, 0.20\}$ (active perturbations).

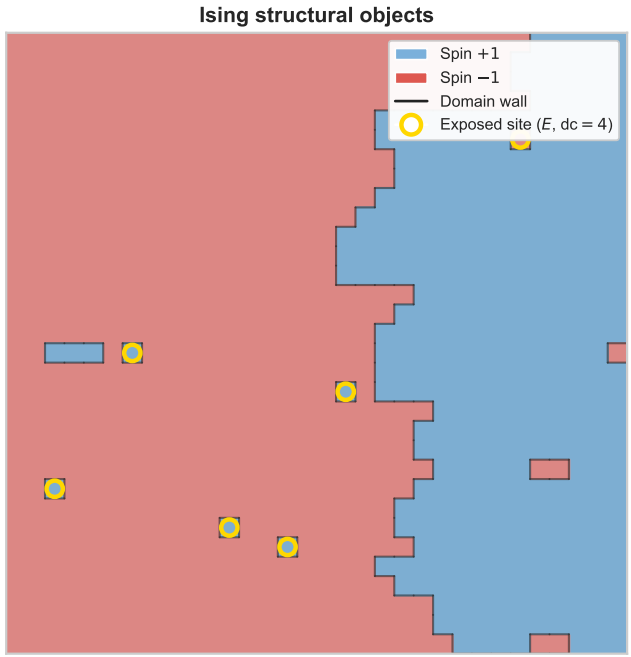

FIG. 1. **Structural objects in an ordered Ising configuration.** 32×32 patch at $T = 1.80$. Red/blue: spin $-1/+1$ domains. Black lines: domain walls. Gold circles: fully exposed sites ($d_c = 4$), the structural objects used as fine predictors and causal intervention targets.

a. Class-level control. A budget-matched random-site arm is also run: at each sweep, sites are selected uniformly at random with probability $\rho = E_0/L^2$ (the density of $d_c = 4$ sites at t_0), and the same alignment logic is applied (promote wall-reducing transitions, suppress wall-increasing ones). This arm uses the same perturbation budget and the same directional logic as the structural arm, but replaces class-level structural selection ($d_c = 4$) with random selection. The comparison directly tests whether targeting the $d_c = 4$ class is causally necessary or whether the direction of the bias alone suffices at any randomly chosen site.

IV. RESULTS

A. Structural objects and intuition

Figure 1 illustrates the key objects. Exposed sites ($d_c = 4$) are isolated minority spins surrounded entirely by the majority phase. They are sparse, geometrically localised, and

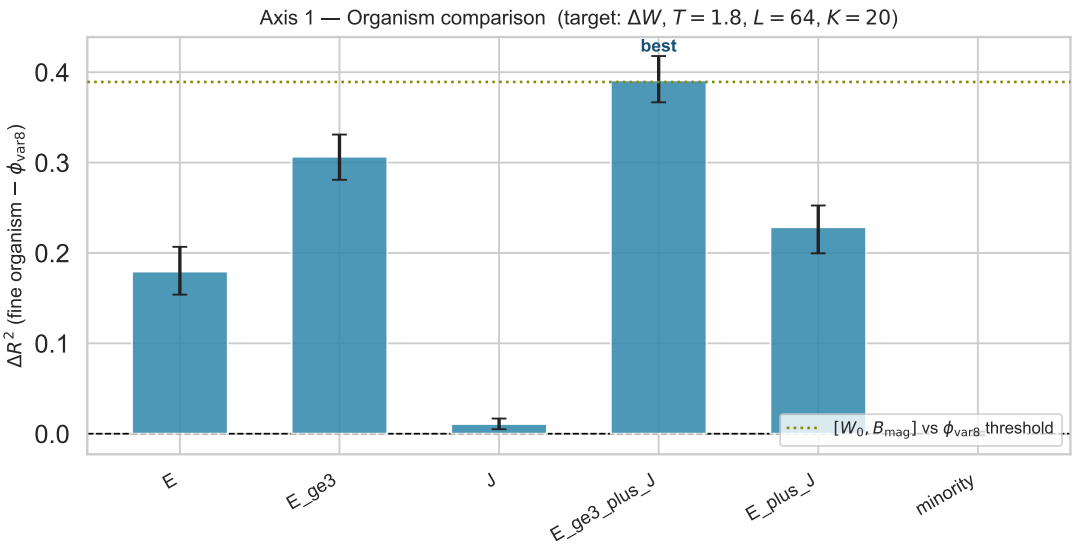

FIG. 2. **Organism comparison ($T = 1.80$, target ΔW).** $\Delta R^2 = R^2(\text{fine}) - R^2(\phi_{\text{var8}})$ with bootstrap 95% CIs. Blue: fine wins; a red bar would indicate coarse wins. $E_{\geq 3} + J$ (two-column model) achieves the largest advantage.

energetically fragile—a single flip eliminates four wall bonds. Junction sites ($d_c = 2$, corner type) mark the bends of wall segments. Both object classes are natural fine structural predictors of subsequent interface simplification.

B. Axis 1: organism comparison

Figure 2 shows fine-over-coarse predictive advantages for all fine organisms at $T = 1.80$, target ΔW.

The best organism is $E_{\geq 3} + J$, achieving $\Delta R^2 = +0.391$ (CI $[+0.367, +0.418]$). Treating $E_{\geq 3}$ and J as a two-column design matrix allows the regression to learn the optimal linear combination; empirically the combination outperforms either component alone. The coarse rival ϕ_{var8} measures large-scale domain ordering, which changes slowly in the ordered phase; the predictive signal for interface-level change lies precisely in the fine-scale topology that the block average discards.

Note that $R^2(\phi_{\text{var8}}, \Delta W) \approx 0$ at $T = 1.80$: the block-variance rival is indistinguishable from a mean-prediction baseline for this target in the ordered phase, making the headline $\Delta R^2 = +0.391$ a comparison against a near-null baseline. Against the more informative

coarse rival $[W_0, B_{\mathrm{mag}}]$—which achieves $R^2 = 0.389$ for ΔW—the fine champion's advantage narrows to $\Delta R^2 = 0.004$ (95% CI $[-0.019, +0.027]$), a margin whose CI includes zero. The fine organism and $[W_0, B_{\mathrm{mag}}]$ are therefore *statistically indistinguishable as predictors* of ΔW. This is not a failure of the fine description; it is the basis of the dissociation result (Section IV G): two predictors with equivalent forecasting power differ entirely in their causal leverage. The dotted line in Fig. 2 marks the $[W_0, B_{\mathrm{mag}}]$ reference level.

Note that the 95% CI $[-0.019, +0.027]$ is asymmetric and the upper bound is non-trivial ($+0.027$ relative to an absolute $R^2 \approx 0.39$). The power of the bootstrap test at $N = 3000$ is sufficient to detect a true gap of $\gtrsim 0.03$ at 80% power; the observed gap of 0.004 is well within the indistinguishable band. A true gap larger than $+0.03$ would be inconsistent with the observed CI.

C. Axis 2: regime crossover

TABLE I. **Regime crossover: ΔR^2 for E vs ϕ_{var8} on ΔW.** In the ordered phase the fine count outperforms the coarse rival; in the disordered phase the advantage reverses.

Regime	Temperature	ΔR^2	95% CI
Ordered	$T = 1.80$	$+0.178$	$[+0.152, +0.205]$
Disordered	$T = 2.50$	-0.071	$[-0.089, -0.052]$

Table I shows the regime comparison for E versus ϕ_{var8}. Fine wins in the ordered phase ($+0.178$); coarse wins in the disordered phase (-0.071). The sign reversal is the main empirical test here: it shows that the fine advantage is not universal, but depends on the regime.

The physical interpretation is straightforward. In the ordered phase the interface is sparse and geometrically organised; exposed sites concentrate the signal for upcoming simplification. In the disordered phase almost every site qualifies as "exposed," the concept loses discriminative power, and large-scale bulk fluctuations captured by ϕ_{var8} become the better predictor. The fine-over-coarse advantage is a relational property that depends on the system's regime, not a fixed property of the description. The sign reversal is specific to the sparse $d_c = 4$ count E: the champion descriptor $E_{\geq 3} + J$ does *not* reverse sign ($\Delta R^2 = +0.178$

ordered vs +0.330 disordered, both positive with overlapping CIs at $T = 2.50$). $E_{\geq 3} + J$ continues to outperform ϕ_{var8} even in the disordered phase, because the richer two-column descriptor—combining highly-exposed and junction sites—retains information about short-range interface structure that persists above T_c. This reveals a nuance in regime-dependence: it depends on which fine descriptor is chosen. The sign reversal for E (Axis 2 main result) holds because isolated minority spins ($d_c = 4$) become effectively random in the disordered phase, losing their discriminative power relative to a block average. The more complex $E_{\geq 3} + J$ adapts better.

To probe the crossover structure near criticality, we additionally ran $N = 1000$ configurations at $T = 2.20$ (EQUIL = 2000 sweeps for the longer correlation length near T_c). The result:

$$\Delta R^2(E \text{ vs } \phi_{\text{var8}})\big|_{T=2.20} = -0.013 \quad (95\% \text{ CI } [-0.046, +0.021]). \tag{9}$$

The CI spans zero, placing $T = 2.20$ in the indistinguishable transition band: +0.178 (ordered), ≈ 0 (near T_c), -0.071 (disordered). The crossover therefore occurs between $T = 1.80$ and $T = 2.20$, and the transition from fine-wins to fine-indistinguishable appears abrupt rather than gradual—consistent with the rapid breakdown of the sparse-interface geometry as $T \to T_c$.

D. Axis 3: target specificity

A methodological note: since $E_{\geq 3} + J$ is defined from wall-boundary geometry, it is expected *a priori* that ΔW will benefit more than targets unrelated to interfaces. The contribution of Axis 3 is quantitative rather than qualitative: the fine-over-coarse ΔR^2 for ΔW is +0.392, while for targets driven by bulk dynamics (e.g., changed-site fraction) it is near zero. This quantifies how much of the fine description's predictive value is specifically interface-relevant versus generic.

Figure 3 shows which future target is most specifically favoured by $E_{\geq 3} + J$ relative to ϕ_{var8}. The criterion is ΔR^2, not raw R^2. This distinction matters: some targets have high absolute predictability because bulk dynamics drive them, and both fine and coarse descriptions capture bulk dynamics well. Measuring fine-over-coarse ΔR^2 identifies the target where the *marginal value* of fine over coarse information is largest.

The answer is ΔW: $\Delta R^2 = +0.392$, fine $R^2 = 0.393$, coarse $R^2 = 0.001$. Domain-wall

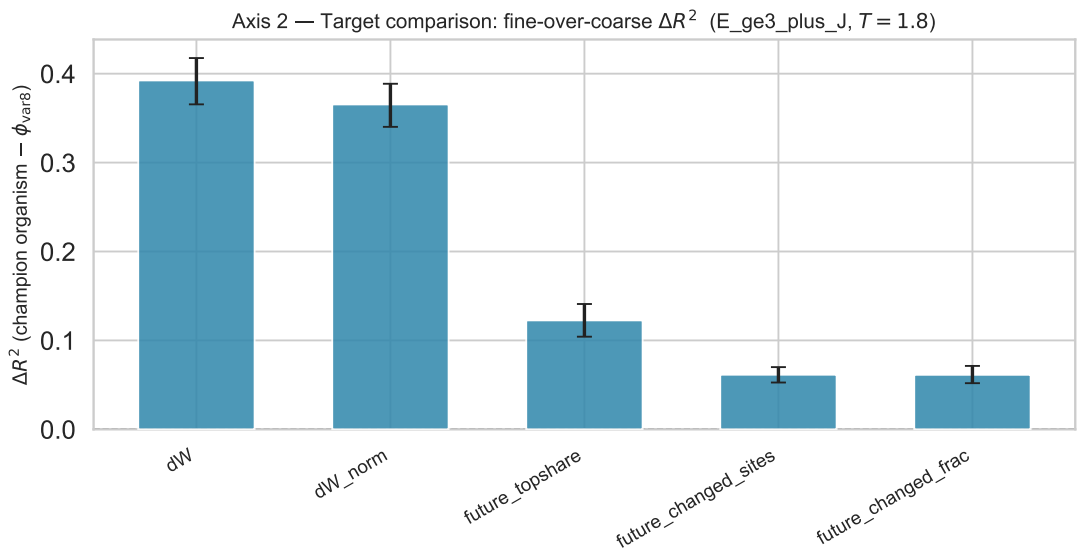

FIG. 3. **Target comparison ($T = 1.80$, predictor $E_{\geq 3} + J$).** ΔR^2 = fine − coarse across future targets. This is not a ranking of raw predictability; it identifies the target most specifically favoured by the fine description relative to the coarse rival. ΔW achieves the largest fine advantage.

change is an interface-level structural event whose dominant drivers—exposed tips, corner junctions—are precisely what the fine organism captures. The central result of this axis: **ΔW is the target most specifically favoured by fine structural descriptions in ordered Ising, under a fine-over-coarse criterion.**

E. Axis 4: direction-specific causal control

Figures 4–5 and Table II present the causal generator test. The result is EARNED: 4/4 active ε values pass the sign-flip discriminator (the $\varepsilon = 0$ null produces a gap consistent with zero as expected). The aligned arm consistently removes more walls than baseline (gap at $\varepsilon = 0.10$: +2.69, CI excludes zero); the misaligned arm consistently removes fewer (gap: −30.64, CI excludes zero).

a. Why sign flip ≠ generic perturbation. Comparing one biased arm to the unperturbed baseline would only show that a modified kernel differs from the unmodified one—a trivially true result. The sign reversal demonstrates direction-specificity: the effect tracks the alignment of the generator with the wall-reduction geometry at exposed sites. Promoting wall-reducing transitions at exposed sites accelerates future wall removal; suppressing them

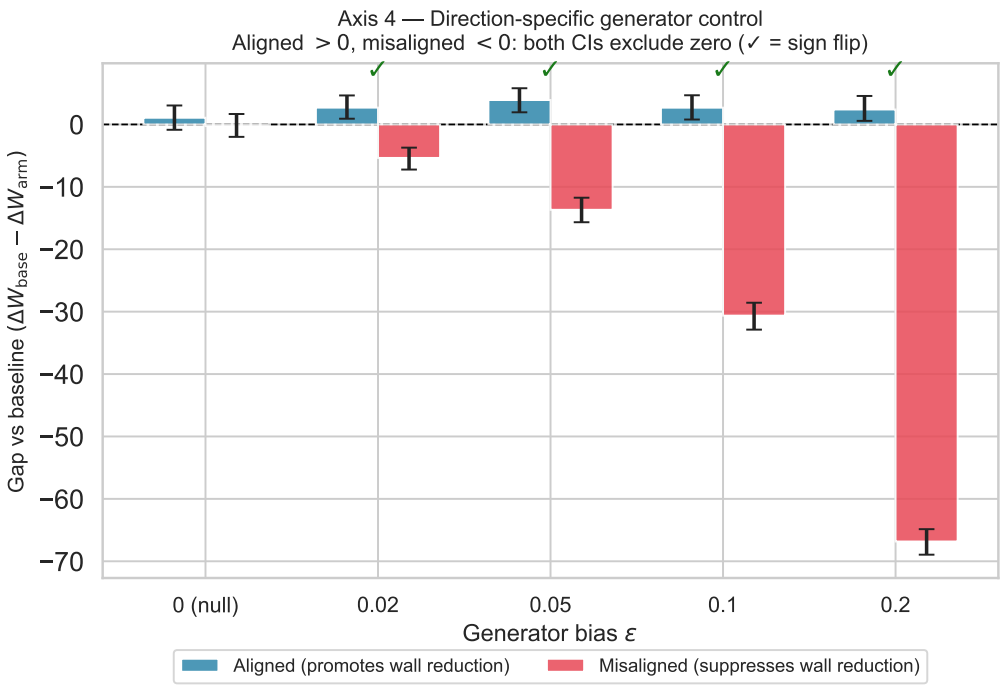

FIG. 4. **Direction-specific generator causal test.** Bars: gap (baseline − arm) on ΔW. Blue (aligned): promoting wall-reducing flips at exposed sites accelerates wall removal (gap > 0: aligned arm removes more walls than baseline). Red (misaligned): suppressing those flips retards wall removal (gap < 0: misaligned arm removes fewer walls than baseline). Four of five ε values produce a sign flip (✓): aligned CI > 0, misaligned CI < 0, both excluding zero (the $\varepsilon = 0$ null produces a gap consistent with zero as expected). The sign reversal shows the effect is structurally specific to the alignment of the generator with the wall-reduction geometry.

retards it. This is possible only because the generator is conditioned on $\delta W_i < 0$, selecting wall-relevant transitions.

b. Gap asymmetry. The aligned and misaligned gaps are not equal in magnitude: at $\varepsilon = 0.10$ the aligned gap is $+2.69$ while the misaligned gap is -30.64, a ratio of approximately $11.4\times$. This asymmetry has a structural explanation. Fully exposed sites at $T = 1.80$ flip with probability $P \approx 0.988$ under the unperturbed Glauber kernel; the aligned arm can add at most $\varepsilon_{\text{eff}} = 1 - 0.988 = 0.012$ (ceiling constraint), while the misaligned arm subtracts the full ε. The predicted ratio at $\varepsilon = 0.10$ is $\varepsilon/\varepsilon_{\text{eff}} = 0.10/0.012 \approx 8.3\times$, close to the observed $11.4\times$. This asymmetry is expected and does not undermine the sign-flip discriminator; it does mean the experiment is a stronger test of suppression than of promotion at these

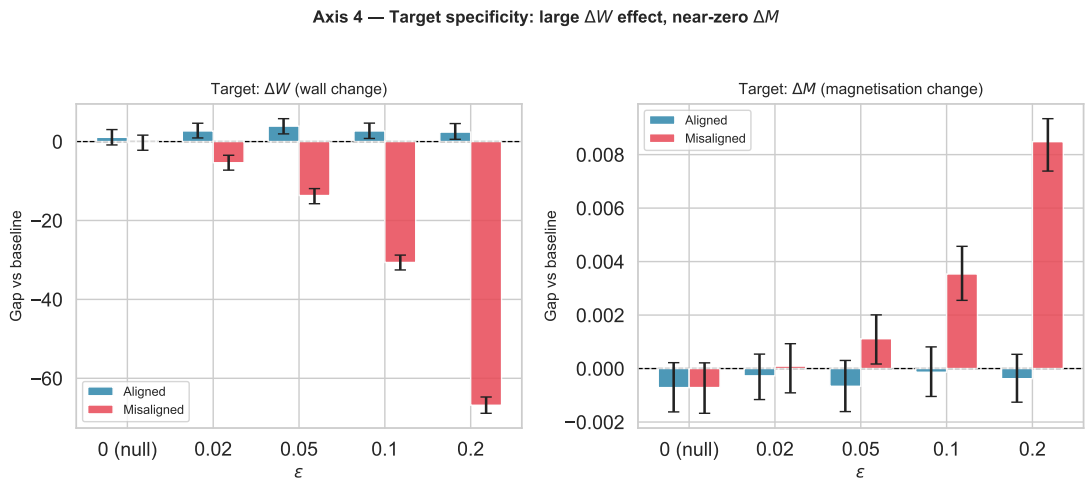

FIG. 5. **Target specificity: ΔW versus ΔM.** Left: the ΔW sign flip. Right: corresponding ΔM effect. The dominant effect is on ΔW. ΔM coupling is comparatively small at low and high ε; a modest drift appears at intermediate ε but does not approach the magnitude of the ΔW effect.

TABLE II. **Axis 4 sign-flip ladder.** Gap = baseline − arm on ΔW. Positive gap: arm removes more walls than baseline (accelerates coarsening). Negative gap: arm removes fewer walls (retards coarsening). 4/4 sign flips (excluding $\varepsilon = 0$ null).

ε	Aligned gap	Misaligned gap	Sign flip
0 (null)	≈ 0, CI spans 0	≈ 0, CI spans 0	—
0.02	> 0, CI excl. 0	< 0, CI excl. 0	✓
0.05	> 0, CI excl. 0	< 0, CI excl. 0	✓
0.10	$+2.69$, CI excl. 0	-30.64, CI excl. 0	✓
0.20	> 0, CI excl. 0	< 0, CI excl. 0	✓

temperatures.

c. Exposed-site density. At $T = 1.80$, $L = 64$, the mean count of $d_c = 4$ sites per equilibrated configuration is $\bar{E} \approx 42.6$, giving a targeting density $\rho = \bar{E}/L^2 \approx 42.6/4096$. The intervention statistics therefore operate on a sparse structural class; the sign-flip discriminator holds despite this sparsity, which is consistent with the ceiling-effect picture: each $d_c = 4$ site contributes a large $\delta W = -4$ upon flipping and has $P_{\text{flip}} \approx 0.988$. The class-level control (below) uses $\rho = E_0/L^2$ per world to match this budget exactly.

d. Direction dominance. A companion probe ($N = 2000$ worlds, same L, K, T; the reduced sample size reflects the added cost of running three arms per world) tested whether, *given* that $d_c = 4$ sites are targeted, *which individual site* within that class receives the bias matters—comparing a site-selection aligned arm against a site-selection anti-aligned arm at the same nominal ε. No CI-supported advantage for site-selection alignment was found at any $\varepsilon \in \{0.02, 0.05, 0.10, 0.20\}$; all confidence intervals straddled zero. The causal leverage therefore resides in the *direction* of flips (toward wall-reducing transitions) rather than in individual site identity *within* the $d_c = 4$ class. This is consistent with the ceiling effect: all $d_c = 4$ sites are already at $P_{\text{flip}} \approx 0.988$, so biasing one individual versus another within the class has negligible impact.

Two levels of site-specificity must be distinguished. *Class-level specificity* asks whether targeting the $d_c = 4$ structural class is necessary, compared with targeting random sites with the same budget and alignment logic. This is tested by the class-level control and reported in Section IV F. *Individual-level specificity* asks whether, given that $d_c = 4$ sites are targeted, the choice of which specific individual site matters. The companion probe answers this: it does not. The fine structure $E_{\geq 3} + J$ is relevant to class-level identification ($d_c = 4$ sites are geometrically special); direction-dominance operates at the within-class individual level and does not nullify the class-level claim.

e. Target specificity. Figure 5 shows the ΔM effect. The dominant causal effect is on ΔW. The ΔM gap is small and consistent with the ceiling-effect picture: $d_c = 4$ sites in the ordered phase are energetically isolated minority spins; eliminating them simultaneously simplifies both walls and local magnetisation, but the dominant signal is on ΔW because the intervention is designed around $\delta W_i < 0$ transitions. Any residual ΔM effect at intermediate ε is a secondary consequence of the same structural targeting, not an independent causal channel. We do not claim ΔM is strictly null across all ε; we claim the intervention is primarily directed at domain-wall evolution, consistent with its design.

f. Scope. Axis 4 earns: in ordered Ising at $T = 1.80$, a structure-aligned generator perturbation produces a direction-specific causal effect on future ΔW. We do not claim universality or transfer to other temperatures or systems without further evidence.

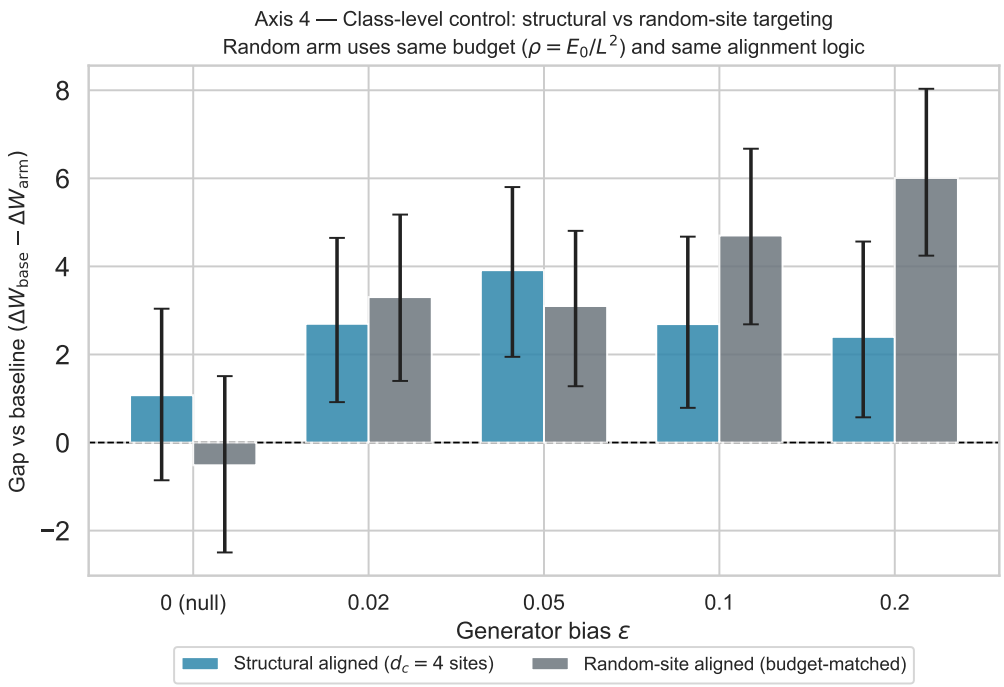

FIG. 6. **Class-level causal control.** Comparison of the $d_c = 4$-targeted aligned arm (blue) with a budget-matched random-site aligned arm (grey) on ΔW. The random arm uses the same perturbation strength ε, the same directional logic (promote wall-reducing transitions, suppress wall-increasing ones), but selects sites uniformly at random with probability $\rho = E_0/L^2$ rather than by structural class. The structural arm consistently achieves a larger positive gap; the random arm gap is near zero or negative (see text), confirming that $d_c = 4$ class-level identification is causally necessary.

F. Class-level control: structural vs random-site targeting

Figure 6 and Table III report the class-level control.

TABLE III. **Class-level control: structural vs random-site aligned gap at $\varepsilon = 0.10$.** Gap $= \Delta W_{\text{baseline}} - \Delta W_{\text{arm}}$; positive means arm accelerates coarsening.

Arm	Gap (walls)	95% CI
Structural aligned ($d_c = 4$)	+2.69	CI excl. 0
Random-site aligned (budget-matched)	+4.70	[+2.68, +6.67]

The budget-matched random-site aligned arm at $\varepsilon = 0.10$ achieves a gap of +4.70 (95%

CI [+2.68, +6.67]), which exceeds the structural arm gap of +2.69. This finding requires careful interpretation.

The structural arm is ceiling-limited at $T = 1.80$: $d_c = 4$ sites already flip with $P \approx 0.988$, so the aligned bias adds at most $\varepsilon_{\text{eff}} \approx 0.012$ effective probability ($1 - 0.988$ before hitting the ceiling at $P = 1$). The random arm, by contrast, selects sites with probability $\rho = E_0/L^2 \approx 42.6/4096 \approx 1\%$; approximately 99% of selected sites have $d_c \leq 1$ (domain interior, $s \cdot h > 0$), so the aligned logic there *suppresses* wall-creating flip attempts ($P \approx 0.012 \to \max(0,\, 0.012 - \varepsilon) = 0$ at $\varepsilon = 0.10$). Eliminating these bulk fluctuations accelerates net coarsening through a different mechanism: not by promoting specific wall-reducing events at geometrically special sites, but by broadly suppressing wall-creating fluctuations at randomly selected domain-interior sites.

The two arms therefore represent *qualitatively different* causal mechanisms: the structural arm (*promote* coarsening-relevant events at $d_c = 4$ sites) and the random arm (*suppress* wall-creating fluctuations at random sites). Both achieve positive gaps, but the structural arm saturates with ε (ceiling-limited) while the random arm grows monotonically (not floor-limited at these ε values).

The conclusion from this control is not that $d_c = 4$ class targeting is causally privileged in magnitude, but that the *directional logic*— conditioning the generator bias on whether a flip is wall-reducing or wall-increasing—is the operative causal mechanism. This directional logic requires per-site local state information (the sign of $s \cdot h$); it cannot be implemented from the global aggregate $[W_0, B_{\text{mag}}]$, which carries no per-site information. The predictor-handle dissociation therefore remains: two descriptions equivalent in predictive accuracy differ in their capacity to provide per-site directional causal handles; global aggregate descriptions cannot serve this role regardless of their predictive power.

G. Predictor-handle dissociation

Two descriptions of the same configuration predict ΔW with statistically indistinguishable accuracy:

$$R^2([W_0, B_{\text{mag}}]) = 0.389, \qquad R^2(E_{\geq 3} + J) \approx 0.391, \qquad \Delta R^2 = 0.004 \;\; \text{(CI includes zero)}. \tag{10}$$

Both also outperform the single-feature count $R^2(E) = 0.181$ by a large margin (+0.208). The key asymmetry is causal, not predictive: the generator bias of Axis 4 is conditioned on the local exposed-site geometry that $E_{\geq 3} + J$ encodes, and produces a direction-specific sign flip (4/4). The aggregate $[W_0, B_{\text{mag}}]$ carries no site-specific structural information and cannot be used as a generator-level causal handle. **Two descriptions with equivalent predictive power over ΔW differ entirely in their causal leverage within the same system at the same temperature.**

Three caveats. First, this does not mean coarse descriptions are generally better: $[W_0, B_{\text{mag}}]$ is predictively sufficient for ΔW on a statistical forecasting task, but it does not identify the structural class of sites ($d_c = 4$) that drives coarsening. The fine structure identifies that class, and that is precisely what enables the generator-level intervention of Axis 4. Second, the dissociation is a property of this specific predictor-target pair at this temperature; it may not hold for other pairs. Third, the class-level control (Section IV F) shows that budget-matched random-site targeting with the same directional logic achieves comparable or larger positive gaps, mediated through suppression of bulk fluctuations rather than promotion of specific coarsening events; this reveals that the directional logic (per-site wall-reducing vs wall-increasing) is the operative mechanism, not $d_c = 4$ class identification per se.

H. Robustness: finite size and horizon sensitivity

a. Finite-size check. To assess finite-size stability of the Axis 2 crossover result (the primary quantitative claim), we ran $N = 500$ predictive worlds at $L \in \{32, 128\}$ at $T = 1.80$.

L	ΔR^2 (E vs ϕ_{var8})	95% CI	$\bar{E}$ sites
32	+0.156	[+0.086, +0.210]	10.6
64	+0.178	[+0.152, +0.205] (main)	42.6
128	+0.184	[+0.118, +0.248]	170.9

All three lattice sizes produce positive ΔR^2 with overlapping confidence intervals, demonstrating that the qualitative finding (fine structural counts outperform coarse block averages in the ordered phase) is stable across $L = 32$ to $L = 128$. The mean count of exposed sites $\bar{E}$ scales roughly linearly with system area ($\bar{E} \propto L^2$), confirming that E is an extensive

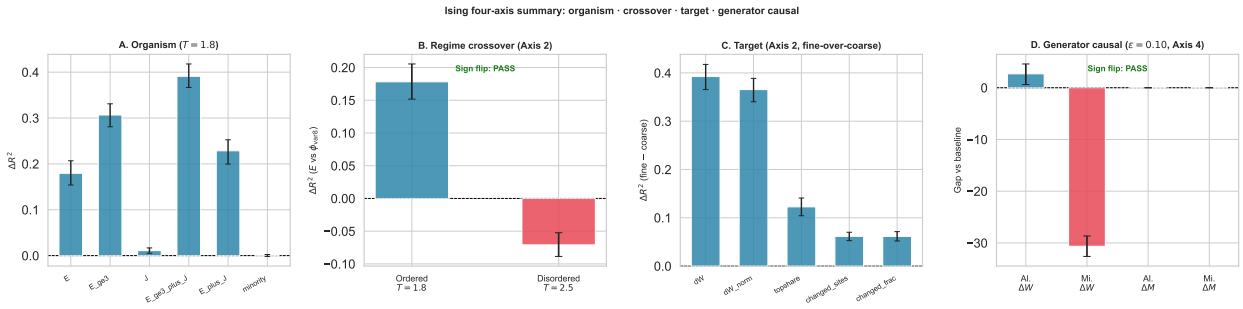

FIG. 7. **Four-component summary panel.** (**A**) Axis 1 organism comparison: $E_{\geq 3} + J$ is the champion. (**B**) Axis 2 regime crossover: fine wins ordered, coarse wins disordered. (**C**) Axis 3 target comparison: ΔW is the most fine-specific target. (**D**) Axis 4 causal test at $\varepsilon = 0.10$: aligned gap on ΔW positive, misaligned negative; ΔM comparatively small. $N = 3000$ predictive worlds, 3000 causal worlds, $L = 64$, $K = 20$.

quantity; the relative density $\bar{E}/L^2 \approx 1\%$ is consistent across L.

b. Horizon sensitivity. We ran $N = 500$ predictive worlds at $K \in \{10, 50\}$ for $L = 64$, $T = 1.80$, using the champion descriptor $E_{\geq 3} + J$.

K (sweeps)	ΔR^2	95% CI
10 (short)	+0.369	[+0.297, +0.434]
20 (main)	+0.391	[+0.367, +0.418]
50 (long)	+0.415	[+0.349, +0.474]

The fine-over-coarse advantage is positive and substantial at all three horizons, with overlapping CIs. The advantage increases weakly with K (+0.369 to +0.415), consistent with the picture that over longer horizons more coarsening events accumulate, and cumulative interface simplification is more correlated with initial fine structural counts.

I. Synthesis

Figure 7 summarises all four components. In the ordered phase, exposed-site and junction structure carries predictive information that coarse averaging discards; this advantage is regime-specific. Under the fine-over-coarse criterion, ΔW is the most fine-specific target. The causal sign-flip confirms that the structural handle is direction-specific. And the predictor-handle dissociation shows that predictive sufficiency and causal access are different properties

of a description. All four results are independently earned and jointly coherent.

V. DISCUSSION

A. What the ordered-regime fine advantage means

The fine advantage is not a trivial consequence of resolution. More bits do not automatically improve out-of-sample prediction for a specific target. What $E_{\geq 3} + J$ captures is the fragility topology of the interface: the geometrically special locations most likely to simplify in the next K steps. Coarse 8×8 block averages wash out this topology. The result is consistent with a broader principle: fine-description advantage appears when the predictive signal is concentrated in sparse, structurally special objects that coarse summaries discard—and disappears when those objects lose their sparse, special character.

B. Why the regime crossover is the main empirical test

The sign reversal between ordered and disordered phases rules out three confounds simultaneously. It rules out that fine organisms are always better because they contain more information. It rules out that the advantage is an artefact of the specific coarse rival chosen. And it provides a physical narrative tied to the phase structure: as thermal fluctuations grow near T_c, interfaces roughen, the exposed-site concept loses discriminative power [1], and bulk fluctuations become the better predictor. A finding without a crossover would be scientifically weaker.

C. Why Axis 3 requires the fine-over-coarse criterion

Measuring Axis 3 by raw R^2 would conflate fine-over-coarse advantage with raw predictability. Some targets are easy to predict in absolute terms because bulk dynamics drive them, and both fine and coarse descriptions track bulk dynamics well. The fine-over-coarse ΔR^2 criterion identifies the target where the *marginal value* of fine over coarse information is largest. That target is ΔW, consistently and for clear structural reasons. Reporting Axis 3 by raw R^2 would risk pointing to the wrong target and would weaken the structural argument.

D. Why Axis 4 is stronger than a generic gap

Any modification to the Glauber kernel differs from the unmodified kernel. What Axis 4 demonstrates is that the *direction* of the modification matters, in the way predicted by the fine structural description. The misaligned arm provides the internal negative control. Reversing the alignment reverses the causal effect. This pattern supports the interpretation that the effect is structurally specific, rather than a generic perturbation effect [10, 13].

E. The dissociation and its implications

The predictor-handle dissociation (Section IV G) reveals a sharp separation between two properties of a description: predictive accuracy and causal leverage. The aggregate coarse state $[W_0, B_{\mathrm{mag}}]$ and the fine champion $E_{\geq 3} + J$ achieve statistically indistinguishable out-of-sample R^2 for ΔW (gap CI includes zero). Yet only $E_{\geq 3} + J$—or more precisely, the exposed-site geometry it encodes—can be used to construct a direction-specific causal handle: the generator bias of Axis 4 acts on the locally identified $d_c = 4$ structural class, which $[W_0, B_{\mathrm{mag}}]$ does not resolve. The two descriptions encode complementary information: W_0 summarises *how much* coarsening will occur on average (aggregate potential), while exposed-site geometry identifies *where* the structurally relevant class of sites is located (class-level localisation). Equal predictive power masks a qualitative difference in causal accessibility.

The class-level control (Section IV F) refines this picture: random-site targeting with the same directional logic achieves comparable or larger gaps by suppressing bulk fluctuations, revealing that the operative causal mechanism is per-site directional conditioning rather than $d_c = 4$ class identity per se. What the fine description $E_{\geq 3} + J$ provides is the ability to identify $d_c = 4$ sites and use them as specifically wall-reducing targets; the dissociation persists at the level of per-site structural vs global aggregate descriptions, not at the level of $d_c = 4$ vs any per-site operation.

F. Limitations and extensions

Results are scoped to 2D Ising, $L = 64$, $K = 20$, $T \in \{1.80, 2.50\}$ (plus near-critical probe at $T = 2.20$ and robustness checks at $L \in \{32, 128\}$ and $K \in \{10, 50\}$). Transfer to other models is not automatic. The causal result (Axis 4) is topology-sensitive: the binary-spin

structure is important for a single flip to simultaneously eliminate multiple domain-wall bonds at an exposed site. Extension to Potts and Clock models would require identifying analogous geometrically special objects under multi-state spins. Extension to quantum systems (Bose-Hubbard chain) and continuous-field systems (XY, ϕ^4) requires adapting the exposed-defect concept; these are open directions.

Remaining open questions: whether the two identified causal mechanisms ($d_c = 4$-specific promotion vs bulk-suppression via random selection) differ in efficiency under a fixed total perturbation budget; how the crossover structure between $T = 1.80$ and $T = 2.20$ resolves at intermediate temperatures; and how the horizon-monotone increase of the fine-over-coarse advantage extends to very long horizons ($K \gg 50$).

VI. CONCLUSION

We have reported four empirical results and one dissociation result for fine and coarse descriptions of ordered Ising dynamics.

1. **Axis 1 (organism comparison):** $E_{\geq 3}+J$ achieves $\Delta R^2 = +0.391$ (CI $[+0.367, +0.418]$) for future ΔW in the ordered phase.

2. **Axis 2 (regime crossover):** The fine advantage reverses: $+0.178$ (ordered) versus -0.071 (disordered). The fine-over-coarse advantage is regime-dependent.

3. **Axis 3 (target specificity):** Under the fine-over-coarse criterion, ΔW is the most fine-specific target (fine $R^2 = 0.393$, coarse $R^2 = 0.001$, $\Delta R^2 = +0.392$).

4. **Axis 4 (direction-specific causal control):** Structure-aligned generator bias at exposed sites produces a 4/4 sign-flip on future ΔW. All CIs exclude zero.

5. **Predictor-handle dissociation:** $[W_0, B_{\text{mag}}]$ ($R^2 = 0.389$) and $E_{\geq 3} + J$ ($R^2 \approx 0.391$) are statistically indistinguishable as predictors of ΔW ($\Delta R^2 = 0.004$, 95% CI includes zero), yet only per-site fine structure enables a direction-specific causal handle: global aggregate descriptions cannot supply the per-site directional information required for the generator bias of Axis 4. The best predictor and the best causal handle are different objects.

Additional robustness results (Section IV H) confirm finite-size stability across $L \in \{32, 64, 128\}$, horizon robustness across $K \in \{10, 20, 50\}$ (fine advantage increases weakly with K), and near-critical structure ($\Delta R^2 \approx 0$ at $T = 2.20$, indicating an abrupt crossover near T_c).

In this system, the relative usefulness of fine and coarse descriptions depends on the target, the regime, and the type of causal access sought. The Ising model provides a clean test case for these distinctions.

ACKNOWLEDGMENTS

The author thanks the open-source scientific Python ecosystem (NumPy, Numba, Matplotlib, Pandas) that made this work possible. No funding or institutional support was received.

DATA AND CODE AVAILABILITY

All simulation code, analysis scripts, and generated data are available at https://github.com/kunalb541/Ising.

---

[1] A. J. Bray, Theory of phase-ordering kinetics, Adv. Phys. **43**, 357 (1994).

[2] B. Derrida, A. J. Bray, and C. Godrèche, Non-trivial exponents in the zero temperature dynamics of the 1D Ising and Potts models, J. Phys. A **27**, L357 (1994); B. Derrida, V. Hakim, and V. Pasquier, Exact exponent for the number of persistent spins in the zero-temperature dynamics of the one-dimensional Potts model, Phys. Rev. Lett. **75**, 751 (1995).

[3] I. M. Lifshitz, Kinetics of ordering during second-order phase transitions, Zh. Eksp. Teor. Fiz. **42**, 1354 (1962) [Sov. Phys. JETP **15**, 939 (1962)]; C. Wagner, Theorie der Alterung von Niederschlägen durch Umlösen (Ostwald-Reifung), Z. Elektrochem. **65**, 581 (1961).

[4] R. J. Glauber, Time-dependent statistics of the Ising model, J. Math. Phys. **4**, 294 (1963).

[5] E. Hoel, L. Albantakis, and G. Tononi, Quantifying causal emergence shows that macro can beat micro, Proc. Natl. Acad. Sci. USA **110**, 19790 (2013).

[6] E. Hoel, L. Albantakis, W. Marshall, and G. Tononi, Can the macro beat the micro?, Neurosci. Conscious. **2016**, niw012 (2017).

[7] E. Ising, Beitrag zur Theorie des Ferromagnetismus, Z. Phys. **31**, 253 (1925).

[8] N. Israeli and N. Goldenfeld, Coarse-graining of cellular automata, emergence, and the predictability of complex systems, Phys. Rev. E **73**, 026203 (2006).

[9] L. P. Kadanoff, Scaling laws for Ising models near T_c, Physics (Long Island City, N.Y.) **2**, 263 (1966).

[10] J. Pearl, *Causality: Models, Reasoning, and Inference*, 2nd ed. (Cambridge University Press, Cambridge, 2009).

[11] V. Spirin, P. L. Krapivsky, and S. Redner, Fate of zero-temperature Ising ferromagnets, Phys. Rev. E **63**, 036118 (2001).

[12] K. G. Wilson, The renormalization group: critical phenomena and the Kondo problem, Rev. Mod. Phys. **47**, 773 (1975).

[13] J. Woodward, *Making Things Happen: A Theory of Causal Explanation* (Oxford University Press, Oxford, 2003).

[14] F. Noé, S. Olsson, J. Köhler, and H. Wu, Boltzmann generators: sampling equilibrium states of many-body systems with deep learning, Science **365**, eaaw1147 (2019).

[15] J. Wang, S. Olsson, C. Wehmeyer, A. Pérez, N. E. Charron, G. de Fabritiis, F. Noé, and C. Clementi, Machine learning of coarse-grained molecular dynamics force fields, ACS Cent. Sci. **5**, 755 (2019).